

In [81]:

```
#!/usr/bin/env python
# coding: utf-8
```

In [82]:

```
# Import dataset (csv.file)
import pandas as pd
df = pd.read_csv('online_shoppers_intention.csv')
```

In [83]:

```
df.head()
```

Out[83]:

	Administrative	Administrative_Duration	Informational	Informational_Duration	ProductRelated
0	0	0.0	0.0	0.0	1
1	0	0.0	0.0	0.0	2
2	0	0.0	0.0	0.0	1
3	0	0.0	0.0	0.0	2
4	0	0.0	0.0	0.0	10

In [84]:

```
# Describe column dimension
numberofcolumns = len(df.columns)
print(numberofcolumns)
```

18

In [85]:

```
# Describe row dimension
numberofrows = len(df.index)
print(numberofrows)
```

12330

In [86]:

```
# Print column names  
print(df.columns)
```

```
Index(['Administrative', 'Administrative_Duration', 'Informational',  
      'Informational_Duration', 'ProductRelated', 'ProductRelated_Duration',  
      'BounceRates', 'ExitRates', 'PageValues', 'SpecialDay', 'Month',  
      'OperatingSystems', 'Browser', 'Region', 'TrafficType', 'VisitorType',  
      'Weekend', 'Revenue'],  
      dtype='object')
```

In [87]:

```
# Check data type  
type(df)
```

Out[87]:

```
pandas.core.frame.DataFrame
```

In [88]:

```
import numpy as np
```

In [89]:

Clean, wrangle, and handle missing data

print(df.describe())

	Administrative	Administrative_Duration	Informational	\
count	12330.000000	12330.000000	12202.000000	
mean	2.315166	80.818611	0.503770	
std	3.321784	176.779107	1.270882	
min	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	
50%	1.000000	7.500000	0.000000	
75%	4.000000	93.256250	0.000000	
max	27.000000	3398.750000	24.000000	

	Informational_Duration	ProductRelated	ProductRelated_Duration	\
count	12330.000000	12330.000000	12330.000000	
mean	34.472398	31.731468	1194.746220	
std	140.749294	44.475503	1913.669288	
min	0.000000	0.000000	0.000000	
25%	0.000000	7.000000	184.137500	
50%	0.000000	18.000000	598.936905	
75%	0.000000	38.000000	1464.157214	
max	2549.375000	705.000000	63973.522230	


	BounceRates	ExitRates	PageValues	SpecialDay	\
count	12330.000000	12330.000000	12195.000000	12330.000000	
mean	0.022191	0.043073	5.911196	0.061427	
std	0.048488	0.048597	18.632116	0.198917	
min	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.014286	0.000000	0.000000	
50%	0.003112	0.025156	0.000000	0.000000	
75%	0.016813	0.050000	0.000000	0.000000	
max	0.200000	0.200000	361.763742	1.000000	

	OperatingSystems	Browser	Region	TrafficType
count	12207.000000	12330.000000	12330.000000	12330.000000
mean	2.123618	2.357097	3.147364	4.069586
std	0.911829	1.717277	2.401591	4.025169
min	1.000000	1.000000	1.000000	1.000000
25%	2.000000	2.000000	1.000000	2.000000
50%	2.000000	2.000000	3.000000	2.000000
75%	3.000000	2.000000	4.000000	4.000000
max	8.000000	13.000000	9.000000	20.000000

In [90]:

```
import matplotlib.pyplot as plt
import seaborn as sns

#visualizing outliers with boxplot
plt.figure(figsize = (15,10))
ax = sns.boxplot(data=df)
```



In [91]:

```
# replace outliers with mean Administrative
df['Administrative'].values[df['Administrative'] >= 10] = 2.315166
print(df['Administrative'].describe())
```

```
count    12330.000000
mean         1.837388
std         2.382258
min          0.000000
25%          0.000000
50%          1.000000
75%          3.000000
max          9.000000
Name: Administrative, dtype: float64
```

In [93]:

```
#replace outliers with mean Administrative_Duration
Q3 = 93.256250
UpperBoundary = Q3*1.5

df['Administrative_Duration'].values[df['Administrative_Duration'] >= UpperBoundary] = 80.818611
print(df['Administrative_Duration'].describe())
```

```
count      12330.000000
mean         34.374434
std          40.305343
min           0.000000
25%           0.000000
50%           7.500000
75%          80.818611
max         139.681818
Name: Administrative_Duration, dtype: float64
```

In [94]:

```
# wrangle Informational
print(df['Informational'].describe())
```

```
count      12202.000000
mean         0.503770
std          1.270882
min           0.000000
25%           0.000000
50%           0.000000
75%           0.000000
max          24.000000
Name: Informational, dtype: float64
```

In [95]:

```
df['Informational'].values[df['Informational'] >= 0] = 0
print(df['Informational'].describe())
```

```
count      12202.0
mean         0.0
std          0.0
min          0.0
25%          0.0
50%          0.0
75%          0.0
max          0.0
Name: Informational, dtype: float64
```

In [96]:

```
#fill in missing values with 0  
df['Informational'].fillna(0)
```

Out[96]:

```
0      0.0  
1      0.0  
2      0.0  
3      0.0  
4      0.0  
...  
12325  0.0  
12326  0.0  
12327  0.0  
12328  0.0  
12329  0.0  
Name: Informational, Length: 12330, dtype: float64
```

In [97]:

```
#wrangle Informational_Duration  
df['Informational_Duration'].describe()
```

Out[97]:

```
count    12330.000000  
mean         34.472398  
std        140.749294  
min          0.000000  
25%          0.000000  
50%          0.000000  
75%          0.000000  
max        2549.375000  
Name: Informational_Duration, dtype: float64
```

In [98]:

```
df['Informational_Duration'].values[df['Informational_Duration'] >= 0] = 0  
print(df['Informational_Duration'].describe())
```

```
count    12330.0  
mean         0.0  
std         0.0  
min         0.0  
25%         0.0  
50%         0.0  
75%         0.0  
max         0.0  
Name: Informational_Duration, dtype: float64
```

In [99]:

```
#wrangle Product_Related
Outlier = 38*1.5
df['ProductRelated'].values[df['ProductRelated'] >= Outlier] = 31.731468
print(df['ProductRelated'].describe())
```

```
count    12330.000000
mean       19.772587
std       13.871250
min        0.000000
25%        7.000000
50%       18.000000
75%       31.000000
max       56.000000
Name: ProductRelated, dtype: float64
```

In [100]:

```
#wrangle ProductRelated_Duration
Q3 = 1464.157214
Q1 = 184.1375
Outlier = (Q3-Q1)*1.5
df['ProductRelated_Duration'].values[df['ProductRelated_Duration'] >= Outlier] = 1194.7462
20
print(df['ProductRelated_Duration'])
```

```
0         0.000000
1        64.000000
2         0.000000
3         2.666667
4       627.500000
...
12325    1783.791667
12326     465.750000
12327     184.250000
12328     346.000000
12329      21.250000
Name: ProductRelated_Duration, Length: 12330, dtype: float64
```

In [101]:

```
#describe wrangled data frame
df.describe()
```

Out[101]:

	Administrative	Administrative_Duration	Informational	Informational_Duration	ProductRela
count	12330.000000	12330.000000	12202.0	12330.0	12330.0000
mean	1.837388	34.374434	0.0	0.0	19.7720
std	2.382258	40.305343	0.0	0.0	13.8710
min	0.000000	0.000000	0.0	0.0	0.0000
25%	0.000000	0.000000	0.0	0.0	7.0000
50%	1.000000	7.500000	0.0	0.0	18.0000
75%	3.000000	80.818611	0.0	0.0	31.0000
max	9.000000	139.681818	0.0	0.0	56.0000

In [102]:

```
df.to_csv('finalproject.csv')
```

In [103]:

```
import seaborn as sns
import statsmodels.api as sm
import pymysql.cursors
import pandas as pd
import matplotlib
import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import axes3d
import numpy as np
import os
```

In [104]:

```
histogram = list(df.describe())
```

In [105]:


```
corr = df.corr().round(2)
```

In [106]:

```
mask = np.triu(np.ones_like(corr, dtype = bool))
```


In [107]:

```
f, ax = plt.subplots(figsize = (12,8))
```



In [108]:


```
cmap = sns.diverging_palette(230,20,as_cmap = True)
```

In [109]:

```
sns.heatmap(corr, mask = mask, annot = True, center =0, linewidths = 1, cmap = cmap)
```

Out[109]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x2428e9c9bb0>
```



In [110]:

```
((df['Month'].groupby([df['VisitorType'], df['Revenue'], df['SpecialDay']]).count())/(df['VisitorType'].count())).round(4)*100
```

Out[110]:

VisitorType	Revenue	SpecialDay	
New_Visitor	False	0.0	9.98
		0.2	0.03
		0.4	0.05
		0.6	0.15
		0.8	0.06
		1.0	0.05
	True	0.0	3.29
		0.2	0.04
		0.4	0.03
		0.6	0.03
Other	False	0.0	0.56
	True	0.0	0.13
Returning_Visitor	False	0.0	64.47
		0.2	1.30
		0.4	1.82
		0.6	2.47
		0.8	2.48
		1.0	1.12
	True	0.0	11.43
		0.2	0.07
		0.4	0.07
		0.6	0.20
		0.8	0.08
		1.0	0.06

Name: Month, dtype: float64

In [111]:

```
df.drop(df[df['ProductRelated_Duration'] > 10000].index, inplace = True)
```

In [112]:

```
(df[['Administrative_Duration','Informational_Duration','ProductRelated_Duration']].groupby([df['Weekend'], df['Revenue']]).mean()/60).round(2)
```

Out[112]:

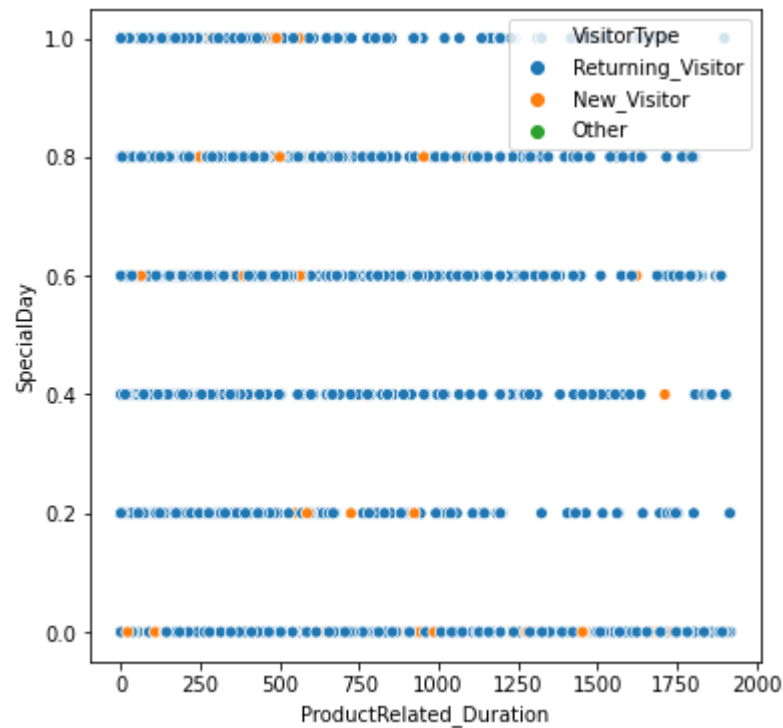
		Administrative_Duration	Informational_Duration	ProductRelated_Duration
Weekend	Revenue			
False	False	0.51	0.0	10.52
	True	0.80	0.0	15.29
True	False	0.59	0.0	11.11
	True	0.80	0.0	15.36

In [113]:

```
sns.scatterplot(data = df, x = "ProductRelated_Duration", y = "SpecialDay", hue = "VisitorType")
```

Out[113]:

<matplotlib.axes._subplots.AxesSubplot at 0x2428eb41340>



In [114]:

```
import numpy as np
from sklearn.preprocessing import LabelEncoder
```

In [115]:

```
le = LabelEncoder()
df['Revenue'] = le.fit_transform(df['Revenue'])
df['Month'] = le.fit_transform(df['Month'])
df['Weekend'] = le.fit_transform(df['Weekend'])
df['VisitorType'] = le.fit_transform(df['VisitorType'])
#f['Weekend'].value_counts()
```

In [60]:

```
y=df.Revenue
x=df.drop('Revenue',axis=1)
```

In [61]:

```
### 70% training data set
### 30% Test data set
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3, random_state = 0)
x_train.head()
```

Out[61]:

	Administrative	Administrative_Duration	Informational	Informational_Duration	ProductRela
11332	1	7.125000	0.0	0.0	
12071	0	0.000000	0.0	0.0	
10023	0	0.000000	0.0	0.0	
6771	9	80.818611	0.0	0.0	
4283	0	0.000000	0.0	0.0	

In [62]:

```
print('Training Features Shape:', x_train.shape)
print('Training Labels Shape:', y_train.shape)
print('Testing Features Shape:', x_test.shape)
print('Testing Labels Shape:', y_test.shape)
```

```
Training Features Shape: (8631, 17)
Training Labels Shape: (8631,)
Testing Features Shape: (3699, 17)
Testing Labels Shape: (3699,)
```

In [74]:

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
import matplotlib.pyplot as plt
import seaborn as sns
model = RandomForestClassifier()
```

In [75]:

```
#df1 = x_train[['Informational','ExitRates', 'PageValues']]
x_train = x_train[['Administrative','Administrative_Duration','Informational_Duration','Pr
oductRelated','ProductRelated_Duration','BounceRates','SpecialDay','Month','Weekend',
                  'Browser','Region','TrafficType','VisitorType','Weekend']]
x_test = x_test[['Administrative','Administrative_Duration','Informational_Duration','Prod
uctRelated','ProductRelated_Duration',
                 'BounceRates','SpecialDay','Month','Weekend','Browser','Region','TrafficT
ype','VisitorType','Weekend']]
model.fit(x_train, y_train)
y_pred = model.predict(x_test)
```

In [76]:

```
# evaluating the model
print("Training Accuracy :", model.score(x_train, y_train))
print("Testing Accuracy :", model.score(x_test, y_test))
```

Training Accuracy : 0.9996524157108099


Testing Accuracy : 0.8313057583130575

In [80]:

```
# confusion matrix
cm = confusion_matrix(y_test, y_pred)
plt.rcParams['figure.figsize'] = (6, 6)
sns.heatmap(cm ,annot = True)
```

Out[80]:

<matplotlib.axes._subplots.AxesSubplot at 0x2428d9194c0>



In [69]:

```
# classification report  
cr = classification_report(y_test, y_pred)  
print(cr)
```

	precision	recall	f1-score	support
0	0.84	0.98	0.91	3077
1	0.49	0.10	0.17	622
accuracy			0.83	3699
macro avg	0.67	0.54	0.54	3699
weighted avg	0.78	0.83	0.78	3699